

# Federated On-Device Training on Arduino Nano 33 BLE

Haoran Li Jiuen Feng Yuhe Bian Zhehao Chen

December 2025

# Outline

- 1 Motivation & Task
- 2 Data & Feature Extraction
- 3 On-Device Neural Network
- 4 Federated Learning Design
- 5 Discussion & Future Work
- 6 Summary

# Project Context

- Hardware: **Arduino Nano 33 BLE + TinyML Shield**.
- Sensors: **IMU** – accelerometer, gyroscope, orientation.
- Target task: classify IMU segments into multiple gestures / movements in “Spell Game”.
- Federated learning on two boards, communication through BLE
- Constraint: training **directly on the board** with very limited RAM and compute.

# Why Federated Learning on Tiny Devices?


- In many IoT scenarios, there are **multiple devices** observing related data distributions.
- Sending raw sensor data to a server is:
  - Bandwidth expensive.
  - Potentially privacy sensitive.
- Federated learning idea:
  - Each device trains a **local model** on its own data.
  - Devices only exchange **model parameters**, not raw data.

# Raw Data from Edge Impulse

- Data collected with Edge Impulse pipeline.
- Stored as `.cbor` messages:
  - Payload includes: `interval_ms`, `sensors`, `values`.
  - Each sample: around **2000 ms**, sampled at **100 Hz**.
- Shape of values:

$$(T, 9) \approx (201, 9),$$

where 9 channels are (accel, gyro, orientation).



# Why Not Use Raw $201 \times 9$ Data on Board?

- Flattened vector length:  $201 \times 9 = 1809$  features.
- With 32-bit floats, a single sample already takes:

$$1809 \times 4 \approx 7.1 \text{ kB}.$$

- On-device training uses:
  - Training data buffer.
  - Network weights and gradients.
  - Stacks, BLE buffers, etc.
- Conclusion: **too large** for robust on-board training.
- Solution: **compress each segment into 75 scalar features.**

# 75-D Feature Extraction (Python)

- For each segment  $v \in \mathbb{R}^{T \times 9}$ :
  - Global statistics per channel (9 dims):
    - mean, std, min, max  $\Rightarrow 9 \times 4 = 36$ .
  - Split time into 3 segments:
    - mean in each segment  $\Rightarrow 3 \times 9 = 27$ .
  - Per-channel energy:

$$\text{energy} = \frac{1}{T} \sum_t v_{t,c}^2 \Rightarrow 9 \text{ dims.}$$

- Magnitude RMS for 3 groups (accel / gyro / orientation):

$$\text{RMS}_{\text{acc}}, \text{RMS}_{\text{gyro}}, \text{RMS}_{\text{ori}} \Rightarrow 3 \text{ dims.}$$

- Total:  $36 + 27 + 9 + 3 = 75$  features per sample.

# Train / Val / Test Split

- Python preprocessing:
  - Read all training and testing .cbor files.
  - Convert to  $X \in \mathbb{R}^{N \times 75}$ , label vector  $y$ .
  - Map label names (e.g., circle, noise) to integers.
  - Randomly shuffle and split into:
    - Train.
    - Validation.
    - Test.
- Store as .npz for PC training and also export as data.h for Arduino.

# Network Architecture on Nano 33 BLE

- Simple fully connected network:

$75 \rightarrow 64 \rightarrow \text{classes\_cnt.}$

- Activation:
  - Hidden layer: ReLU.
  - Output layer: softmax.
- Loss:
  - Cross-entropy between predicted probabilities and one-hot labels.
- Implemented in C with:
  - Manually allocated layers and neurons.
  - Forward and backward propagation using SGD.

# On-Device Data & Normalization

- `data.h` contains:
  - `train_data[train_cnt][75]`.
  - `validation_data[val_cnt][75]`.
  - `test_data[test_cnt][75]`.
  - `train_labels, validation_labels, test_labels`.
  - Feature-wise `feature_min[75], feature_max[75]` from train set.
- Before feeding to the network (on board):

$$x'_j = \frac{x_j - \text{feature\_min}_j}{\text{feature\_max}_j - \text{feature\_min}_j}.$$

- Normalization is implemented as a small C function operating on the `input []` buffer.

# Baseline: Local On-Device Training

- Optimizer: **SGD** with fixed learning rate.
- Example parameters:
  - `#define LEARNING_RATE 0.0015`
  - `#define EPOCH 50`
- Training loop on board:
  - Shuffle training indices.
  - For each sample:
    - 1 Load feature vector into `input []`.
    - 2 Normalize using `feature_min/max`.
    - 3 Forward propagation.
    - 4 Backward propagation, update weights.
- Accuracy monitored on train / val / test after each epoch.

# From Single Board to Federated Setup

- We have two Nano 33 BLE boards:
  - Each board has its own local dataset from different users' gesture recordings.
- Idea: In each round
  - Each board trains locally for several epochs.
  - Boards use BLE to **exchange network weights**.
  - Then **average** the weights as a simple federated aggregation.
- No raw sensor data is transmitted.

# Parameter Packing for BLE

- In the C code, the network is represented as:
  - Layers  $L[i]$ .
  - Neurons with weight vector  $W[]$  and bias  $B$ .
- Function `packUnpackVector(Type)`:
  - PACK: serialize all weights and biases into a flat `WeightBiasPtr[]` buffer.
  - UNPACK: read back from buffer to local network.
  - AVERAGE: average between received buffer and local network and update both.
- This buffer is then sent via BLE characteristic(s) between the two boards.

# Simple Federated Averaging

- In standard FedAvg:

$$w^{(t+1)} = \sum_k \frac{n_k}{N} w_k^{(t+1)},$$

where  $w_k$  is client  $k$ 's local weights.

- In our prototype with two boards:
  - Use equal weighting as an approximation:

$$w_{\text{new}} = \frac{w_A + w_B}{2}.$$

- Implemented in `packUnpackVector(AVERAGE)`.
- Federated round:
  - ① Each board trains locally for some epochs.
  - ② Exchange weight buffers via BLE.
  - ③ Average and update local networks.

# Federated Training Schedule

- Example schedule:
  - Local epochs per round:  $E$  (e.g., 1–5).
  - Number of federated rounds:  $R$ .
- On each board:
  - 1 Repeat local training for  $E$  epochs.
  - 2 Pack parameters and send to peer via BLE.
  - 3 Receive peer parameters into buffer.
  - 4 Call `packUnpackVector(AVERAGE)` to update network.
- Trade-off:
  - Larger  $E$ : fewer communications, more local drift.
  - Smaller  $E$ : more communication, better synchronization.


# PC-Side Reference Training (PyTorch)

- Use exactly the same 75-D feature vectors as on the Arduino:
  - Dataset merged from all .npz files: about 180 training, 45 validation and 58 test samples, 4 classes.
  - Network on PC is identical to on-device network:

$$75 \rightarrow 64 \rightarrow \text{classes\_cnt}$$

- Loss: cross-entropy, optimizer: SGD with  $\text{lr} = 0.0015$ , batch size 1 (to mimic on-device training).
- Results:
  - After about 50 epochs, validation accuracy reaches  $\approx 95\%$ .
  - This PC model is used as an **upper-bound reference** for on-device learning performance.


## Training / Validation Curves



*PC-side centralized training reaches  $\sim 95\%$  val. accuracy in  $\sim 50$  epochs.*

# Federated Training on PC (MPI + FedAvg)

- Federated set-up (FedAvg):
  - Each MPI rank  $r = 1, \dots, K$  loads its local `client_r.npz` and trains a local model for `LOCAL_EPOCHS` with SGD.
  - Rank 0 acts as the server:
    - broadcasts global weights to all clients,
    - gathers updated weights,
    - performs weighted averaging (FedAvg),
    - evaluates on a held-out validation set and logs accuracy.
- Hyper-parameters:  
`NUM_ROUNDS = 50, LOCAL_EPOCHS = 1, BATCH_SIZE = 1, LR = 0.0015.`



# Limitations

- **Optimizer:** on-device training uses plain SGD; no Adam or momentum due to:
  - Code complexity.
  - Extra memory for moment estimates.
- **Communication:**
  - BLE bandwidth is limited.
  - Transmitting full parameter vector can be slow.
- **Scalability:**
  - Prototype only tested with two boards.
  - No central server; logic is peer-to-peer.

# Possible Improvements

- On-device optimization:
  - Implement lightweight variants of Adam / momentum, or adaptive learning rate schedules.
- Model compression:
  - Quantize weights / send only deltas to reduce BLE traffic.
- Protocol:
  - Generalize from 2 boards to  $K$  devices, with either a central aggregator or a gossip protocol.
- Data:
  - Collect more user-specific IMU trajectories to better demonstrate personalization + federated averaging.

# Summary

- Built a **full pipeline**:
  - Raw IMU segments  $\Rightarrow$  75-D features.
  - Python preprocessing and PC training.
  - Exported data.h for on-device training.
- Implemented a small **on-device neural network**:
  - 75–64-classes\_cnt, ReLU + softmax, SGD training.
- Designed and tested a **federated learning prototype**:
  - Two Nano 33 BLE boards exchanging weights via BLE.
  - Simple parameter averaging after local training epochs.
- Demonstrated that even with tight resource constraints, federated ideas can be prototyped on microcontrollers.

# Thank you for listening!