

# Federated On-Device Training on Arduino Nano 33 BLE

Haoran Li Jiuen Feng Yuhe Bian Zhehao Chen

December 2025

# Outline

- 1 Motivation & Task
- 2 Data & Feature Extraction
- 3 On-Device Neural Network
- 4 Federated Learning Design
- 5 Summary

# Project Context

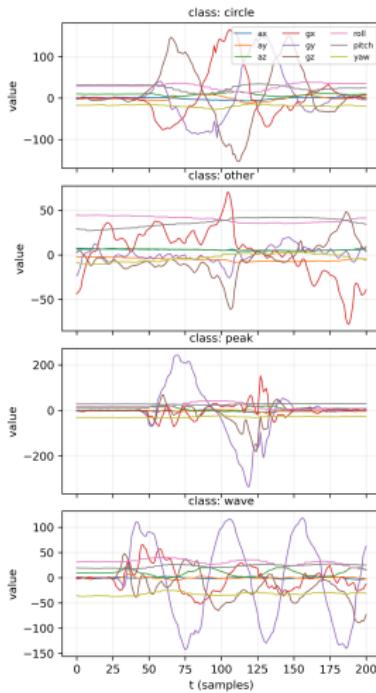
- Hardware: **Arduino Nano 33 BLE + TinyML Shield**.
- Sensors: **IMU** – accelerometer, gyroscope, orientation.
- Target task: classify IMU segments into multiple gestures / movements in “Spell Game”.
- Federated learning on two boards, communication through BLE
- Constraint: training **directly on the board** with very limited RAM and compute. BLE transmitting limitation.

# Raw Data from Edge Impulse

- Data collected with Edge Impulse pipeline.
- Stored as .cbor messages:
  - Payload includes: `interval_ms`, `sensors`, `values`.
  - Each sample: around **2000 ms**, sampled at **100 Hz**.
- Shape of values:

$$(T, 9) \approx (201, 9),$$

where 9 channels are (accel, gyro, orientation).



# 75-D Feature Extraction (Python)

- For each segment  $v \in \mathbb{R}^{T \times 9}$ :
  - Global statistics per channel (9 dims):
    - mean, std, min, max  $\Rightarrow 9 \times 4 = 36$ .
  - Split time into 3 segments:
    - mean in each segment  $\Rightarrow 3 \times 9 = 27$ .
  - Per-channel energy:

$$\text{energy} = \frac{1}{T} \sum_t v_{t,c}^2 \Rightarrow 9 \text{ dims.}$$

- Magnitude RMS for 3 groups (accel / gyro / orientation):

$$\text{RMS}_{\text{acc}}, \text{RMS}_{\text{gyro}}, \text{RMS}_{\text{ori}} \Rightarrow 3 \text{ dims.}$$

- Total:  $36 + 27 + 9 + 3 = 75$  features per sample.

# Network Architecture on Nano 33 BLE

- Simple fully connected network:

$75 \rightarrow 64 \rightarrow \text{classes\_cnt.}$

- Activation:
  - Hidden layer: ReLU.
  - Output layer: softmax.
- Loss:
  - Cross-entropy between predicted probabilities and one-hot labels.
- Implemented in C with:
  - Manually allocated layers and neurons.
  - Forward and backward propagation using SGD.

# From Single Board to Federated Setup

- We have two Nano 33 BLE boards:
  - Each board has its own local dataset from different users' gesture recordings.
- Idea: In each round
  - Each board trains locally for several epochs.
  - Boards use BLE to **exchange network weights**.
  - Then **average** the weights as a simple federated aggregation.
- **Only parameters are sent over BLE.** No raw sensor data is transmitted.

# BLE Communication

## Board A: Leader (BLE Peripheral)

- Advertises BLE service MLLeader and accepts a connection.
- **Sends** the current **global weights** over BLE (split into batches).
- **Receives** updated weights from Board B over BLE.
- Aggregates parameters and starts the next round.

## Board B: Worker (BLE Central)

- Scans for MLLeader, connects, and subscribes to updates.
- **Receives** global weights over BLE and loads them into the NN.
- Runs **local on-device training** on its own IMU dataset.
- **Uploads** its updated weights back to Board A over BLE.

# Simple Federated Averaging

- In standard FedAvg:

$$w^{(t+1)} = \sum_k \frac{n_k}{N} w_k^{(t+1)},$$

where  $w_k$  is client  $k$ 's local weights.

- In our prototype with two boards: Use equal weighting as an approximation:

$$w_{\text{new}} = \frac{w_A + w_B}{2}.$$

# On-Device Training Curves




Local Train with data\_byh




Local Train with data\_fje

- Federated training improves **cross-user generalization** without sharing raw IMU data.

# On-Device Training Curves




TODO



TODO


- Federated training improves **cross-user generalization** without sharing raw IMU data.

# Confusion Matrices: Local Training one Board




Train byh → Test byh

acc = 82.50%



Train byh → Test czh

acc = 58.67%






Train byh → Test fje

acc = 52.83%

- Best performance on **in-domain** test (A→A).
- Accuracy drops on B/C due to **non-IID user motion patterns**.

# Confusion Matrices: Federated Training (fje+byh, FedAvg)



FedAvg → Test byh

FedAvg → Test fje

FedAvg → Test czh

- Federated model reduces cross-user confusion compared to local models.
- Gains come with communication cost (BLE weight exchange each round).


# PC-Side Reference Training (PyTorch)

- Use exactly the same 75-D feature vectors as on the Arduino:
  - Dataset merged from all .npz files: about 180 training, 45 validation and 58 test samples, 4 classes.
  - Network on PC is identical to on-device network:

$$75 \rightarrow 64 \rightarrow \text{classes\_cnt}$$

- Loss: cross-entropy, optimizer: SGD with  $\text{lr} = 0.0015$ , batch size 1 (to mimic on-device training).
- Results:
  - After about 50 epochs, validation accuracy reaches  $\approx 95\%$ .
  - This PC model is used as an **upper-bound reference** for on-device learning performance.

## Training / Validation Curves



*PC-side centralized training reaches  $\sim 95\%$  val. accuracy in  $\sim 50$  epochs.*


# Federated Training on PC (MPI + FedAvg)

- Federated set-up (FedAvg):

- Each MPI rank  $r = 1, \dots, K$  loads its local `client_r.npz` and trains a local model for `LOCAL_EPOCHS` with SGD.
- Rank 0 acts as the server:
  - broadcasts global weights to all clients,
  - gathers updated weights,
  - performs weighted averaging (FedAvg),
  - evaluates on a held-out validation set and logs accuracy.

- Hyper-parameters:

`NUM_ROUNDS = 50, LOCAL_EPOCHS = 1, BATCH_SIZE = 1, LR = 0.0015.`



# Demo Video

Click to watch demo video

# Summary

- Built a **full pipeline**:
  - Raw IMU segments  $\Rightarrow$  75-D features.
  - Python preprocessing and PC training.
  - Exported data.h for on-device training.
- Implemented a small **on-device neural network**:
  - 75–64-classes\_cnt, ReLU + softmax, SGD training.
- Designed and tested a **federated learning prototype**:
  - Two Nano 33 BLE boards exchanging weights via BLE.
  - Simple parameter averaging after local training epochs.
- Demonstrated that even with tight resource constraints, federated ideas can be prototyped on microcontrollers.

# Thank you for listening!